

Nishad Wajge

1-408-318-5346 | nishad.wajge@gmail.com | linkedin.com/in/nishadwajge | github.com/nishadw | nishad.page | San Jose, CA

EDUCATION

University of Maryland

College Park, MD

BS in Computer Science (Concentration: Machine Learning) | Minors: Business, Statistics

Aug 2023 – Dec 2026

Coursework: Applied Probability & Stats, Artificial Intelligence, Computational Game Theory, Computer Vision, Data Science, Differentiable Programming, Machine Learning, Parallel Computing, Sampling Theory, Statistical Computing

EXPERIENCE

Amazon

Jun 2026 – Sep 2026

Software Engineer Intern | Amazon Lab126

Sunnyvale, CA

- Shipped C++ on-device speaker ID module to **22M+** vehicles, never blocking a response (**0ms** perceived latency)
- Increased true accept rate **80%→97%** at **1%** FAR by gating **1.4%** of short, impostor-prone utterances
- Cut embedding latency **4x** to **77ms** on ARM via INT8 ONNX export, enabling real-time speaker recognition
- Built MLOps loop (**1M+** trial eval harness fed by fleet data) for continuous fine-tuning of SLMs, ASR, TTS
- Synced AES-256-GCM encrypted voiceprints across devices via hub/replica merge, keeping IPC off the hot path

Amazon

Jan 2026 – May 2026

ML Engineering Lead (contract) | Amazon Leo

College Park, MD

- Forecasted G7 space policies **12** months ahead, hitting **91%** backtested accuracy via 3-model time-series ensemble
- Engineered ETL backend to process **33K+** documents from 9-year lakehouse into structured relational ontology
- Improved signal-to-noise ratio **42%** across **10** signals via temporal smoothing, reducing false positives **25%**

General Dynamics Information Technology

Jun 2025 – Aug 2025

Software Engineer Intern

Annapolis Junction, MD

- Deployed 0-to-1 edge vision module detecting vehicles and streaming geospatial telemetry to connected users
- Sustained **95%** out-of-sample accuracy on noisy sensor data via projected gradient descent adversarial training
- Architected thread-safe, concurrent 3-stage quantized pipeline hitting **91ms** TPU latency on live video feeds

General Dynamics Information Technology

Jan 2025 – May 2025

Engineering Project Lead (contract)

College Park, MD

- Led **13** engineers shipping a hybrid graph and vector RAG pipeline over **600K+** records to **100+** analysts
- Raised retrieval accuracy **40%** via custom chunking and named-entity recognition, cutting discovery time **95%**

Internal Revenue Service

Jan 2024 – Dec 2024

Software Engineer Intern

Washington, DC

- Upgraded division-wide Jenkins CI/CD (JDK 11→17) for fraud detection on **22.5M+** monthly tax applications
- Halved build times for **1,000+** engineers by migrating a 20-year-old legacy Java system from Ant to Maven

RESEARCH

Gradient Fidelity in Learned World Models | First Author | In preparation

Jacobian-matched world models that recover true simulator gradients for gradient-based planning

Standardized Difficulty Labels for Profiling LLMs | Undergraduate Researcher | NeurIPS 2024

Continuous difficulty labels for six LLM benchmarks via Item Response Theory and Glicko-2 rating models

Should Professionals Consider Their Adversary's Strategy? | First Author | Comput. Stat. 2024

MDP and two-player stochastic game models on PGA Tour data, bounding opponent-aware gain at 0.12 points/round

TECHNICAL SKILLS

Languages: Python, C, C++, CUDA, Rust, Java, Kotlin, JavaScript/TypeScript, SQL, x86 Assembly

Technologies: PyTorch, TensorFlow/Lite, JAX, ONNX Runtime, Hugging Face, OpenCV, Pandas, React, Node.js

Software: AWS (SageMaker, Bedrock, S3), Docker, Jenkins, Neo4j, Qdrant, PostgreSQL, Maven, Linux, Git

Interests: Competitive Golf, Investing, Poker, Film Analysis, Classical Music